

# Parallel Algorithms



# Last time ...

- ▶ Introduction to Big Data
- ▶ Data Collection
- ▶ Assignment #1

## Questions?



# Today ...

- ▶ New classroom – WEB L114
- ▶ **Reminder:** Assignment #1 questionnaire due today
- ▶ Permission codes
  
- ▶ Intro to Parallel Algorithms
  - ▶ Complexity analysis
  - ▶ Work depth
  - ▶ Communication costs



# Parallel Thinking

THE MOST IMPORTANT GOAL OF TODAY'S LECTURE




# Parallelism & beyond ...



**1 ox:** single core performance



**1024 chickens:** parallelism



**tractor:** better algorithms

*If you were plowing a field, which would you rather use?  
Two strong oxen or 1024 chickens?*

*Seymour Cray*



Consider an array  $A$  with  $n$  elements,

Goal: to compute,


$$x = \sum_{i=1}^n \sqrt{A_i}$$

Machine Model  
Programming Model  
Performance analysis



# Von Neumann architecture

- ▶ Central Processing Unit (CPU, Core)
- ▶ Memory
- ▶ Input/Output (I/O)
- ▶ One instruction per unit/time
- ▶ **Sequential**





# Characterizing algorithm performance

- ▶  $O$ -notation

- ▶ Given an input of size  $n$ , let  $T(n)$  be the total time, and  $S(n)$  the necessary storage
- ▶ Given a problem, is there a way to compute lower bounds on storage and time → **Algorithmic Complexity**

- ▶  $T(n) = O(f(n))$  means

$T(n) \leq cf(n)$  , where  $c$  is some unknown positive constant

compare algorithms by comparing  $f(n)$ .





# Scalability

- ▶ Scale Vertically → scale-up
  - ▶ Add resources to a single node
  - ▶ CPU, memory, disks,
- ▶ Scale Horizontally → scale-out
  - ▶ Add more nodes to the system



# Parallel Performance

- ▶ Speedup

best sequential time/time on p processors

- ▶ Efficiency

speedup/p, ( $< 1$ )

## Scalability



# Amdahl's Law

- ▶ Sequential bottlenecks:

Let  $s$  be the percentage of the overall work that is sequential

- ▶ Then, the speedup is given by

$$S = \frac{1}{s + \frac{1-s}{p}} \leq \frac{1}{s}$$



# Gustafson

Sequential part should be independent of the problem size

Increase problem size, with increasing number of processors



# Strong & Weak Scalability

Increasing number of cores

**Strong** (fixed-sized) scalability

keep problem size fixed

**Weak** (fixed-sized) scalability


keep problem size/core fixed



# Parallel Programming

- ▶ Partition Work      Data & Tasks
- ▶ Determine Communication
- ▶ Agglomeration to number of available processors
- ▶ Map to processors
- ▶ Tune for architecture









# Work/Depth Models

- ▶ Abstract programming model
- ▶ Exposes the parallelism
  - ▶ Compute work  $W$  and depth  $D$
  - ▶  $D$  - longest chain of dependencies
  - ▶  $P = W/D$
- ▶ Directed Acyclic Graphs
- ▶ Concepts
  - ▶ parallel for (data decomposition)
  - ▶ recursion (divide and conquer)



# Work/Depth Models

- ▶ Abstract programming model
- ▶ Exposes the parallelism
  - ▶ Compute work  $W$  and depth  $D$
  - ▶  $D$  - longest chain of dependencies
  - ▶  $P = W/D$
- ▶ Directed Acyclic Graphs
- ▶ Concepts
  - ▶ parallel for (data decomposition)
  - ▶ recursion (divide and conquer)



# Parallel vs sequential for

- ▶ Dependent statements
  - ▶  $W = \sum W_i$
  - ▶  $D = \sum D_i$
- ▶ Independent statements
  - ▶  $W = \sum W_i$
  - ▶  $D = \max(D_i)$



# Parallel Sum – language model

```
// Recursive implementation
```

```
Algorithm SUM(a, n)
```

```
// Input: array a of length  $n = 2^k, k = \log n$ 
```

```
    parallel_for i  $\leftarrow$  [0, n/2)
```

```
        b(i)  $\leftarrow$  a(2i) + a(2i+1)
```

```
    return SUM(b); //  $W\left(\frac{n}{2}\right), D\left(\frac{n}{2}\right)$ 
```

**Complexity:**

$$D(n) = D\left(\frac{n}{2}\right) + O(1) = O(\log n)$$

$$W(n) = W\left(\frac{n}{2}\right) + O(n) = O(n)$$



# Prefix Sums

SCAN



# Prefix sum or a scan

▶  $[4 \ 2 \ 1 \ 8] \rightarrow [4 \ 4 + 2 \ 4 + 2 + 1 \ 4 + 2 + 1 + 8]$   
 $\rightarrow [4 \ 6 \ 7 \ 15]$

▶ Obvious sequential algorithm

$$D_s(n) = O(n)$$

▶ Not parallelizable, need to change algorithm

▶ Remark: algorithm generalizes to any associative operation



# Parallel prefix sum

- ▶ Input:  $\{x_i\}, \quad i = 1, \dots, n$
- ▶ Output:  $\{s_i\}, \quad i = 1, \dots, n \quad | \quad s_i = \sum_{k=1}^{i-1} x_k$
- ▶ Sequential algorithm: best  $T_s(n) = O(n)$
- ▶ Parallel
  - ▶ Work/Depth  $W(n) = O(n), \quad D(n) = O(\log n)$
  - ▶ Recursive and iterative implementation
- ▶ Binary tree (like sum)
  - ▶ Two tree traversals: bottom-up and top-down





# Key Idea

- ▶ Divide & Conquer + Recursion
  - ▶ Break the scan in parts
  - ▶ Compute subcomponents
  - ▶ Combine
- ▶ Implementation
  - ▶ Recursive or Iterative



# Divide & Conquer



# Recursive Scan

```
s = Rec_Scan(a,n)
// s[k]= a[0]+a[1]+...+a[k-1], k=0,...,n-1
1. if n=1 { s[0]←a[0]; return; } //base case
2. parfor (i=0; i<n/2; ++i)
    b[i] ← a[2i] + a[2i+1];
3. c ← Rec_Scan (b, n/2);
4. s[0] ← a[0];
5. parfor (i=0; i<n; ++i)
    if isOdd(i)
        s[i] ← c[i/2];
    elseif isEven(i)
        s[i] ← c[i/2] + a[i];
```

Iterative ?



# Parallel Select

- ▶ Select numbers  $<$  pivot

$A \leftarrow [1 \ 2 \ 3 \ 0 \ 4 \ 0 \ 2 \ 3 \ 0 \ 1 \ 3 \ 4]$

$\text{pivot} \leftarrow 2$

$[1 \ 0 \ 0 \ 1 \ 0 \ 1 \ 0 \ 0 \ 1 \ 1 \ 0 \ 0]$

$[1 \ 1 \ 1 \ 2 \ 2 \ 3 \ 3 \ 3 \ 4 \ 5 \ 5 \ 5]$



# Parallel Select the $a_i < pv$

```
[l,m] ← select_lower (a, n, pv)
// t = t[0,...,n-1]
parfor (i=0; i<n; ++i)    t[i] ← a[i] < pv;
s ← scan (t);  m ← s(n-1);
parfor (i=0; i<n; ++i)    if t[i]  l[s[i] - 1] ← a[i];
```

$$W(n) = O(n)$$
$$D(n) = O(\log n)$$

