

Parallel Algorithms

PART 2



Last time ...

- ▶ Introduction to Parallel Algorithms
- ▶ Complexity analysis
- ▶ Work/Depth model
- ▶ Prefix Sum, Parallel Select

Questions?



Parallel Select

- ▶ Select numbers $<$ pivot

$A \leftarrow [1 \ 2 \ 3 \ 0 \ 4 \ 0 \ 2 \ 3 \ 0 \ 1 \ 3 \ 4]$

$\text{pivot} \leftarrow 2$

$[1 \ 0 \ 0 \ 1 \ 0 \ 1 \ 0 \ 0 \ 1 \ 1 \ 0 \ 0]$

$[1 \ 1 \ 1 \ 2 \ 2 \ 3 \ 3 \ 3 \ 4 \ 5 \ 5 \ 5]$

Parallel Select the $a_i < pv$

```
[l,m] ← select_lower (a, n, pv)
// t = t[0,...,n-1]
parfor (i=0; i<n; ++i)    t[i] ← a[i] < pv;
s ← scan (t);  m ← s(n-1);
parfor (i=0; i<n; ++i)    if t[i]  l[s[i] - 1] ← a[i];
```

$$W(n) = O(n)$$
$$D(n) = O(\log n)$$

Today ...

Intro to Parallel Algorithms

- ▶ Parallel Search
- ▶ Parallel Sorting
 - ▶ Merge sort
 - ▶ Sample sort
 - ▶ Bitonic sort
- ▶ Communication costs



Parallel Search

- ▶ Problem Description
 - ▶ Given a sorted list X of size n and an element y
 - ▶ Find the index $i \mid x_i \leq y < x_{i+1}$
- ▶ Sequential
 - ▶ Use binary search
 - ▶ $O(\log n)$ time
- ▶ Work depth
 - ▶ `parfor(i) if $x_i \leq y < x_{i+1}$ return i; // no duplicates $W = n, D = 1$`
- ▶ PRAM
 - ▶ $O\left(\frac{\log n}{\log p}\right)$ using p processes



Ranking

- ▶ Given ordered lists, A, B of lengths s, t

- ▶ Define:

$$\text{rank}(z: A) \leftarrow \text{number of elements } a_i \mid a_i \leq z$$

- ▶ Define:

$$\begin{aligned} \text{rank}(B: A) &:= (r_1, r_2, \dots, r_t) \\ r_i &\leftarrow \text{rank}(b_i: A) \end{aligned}$$



Ranking

- ▶ $A = [7 \quad 13 \quad 25 \quad 26 \quad 31 \quad 54]$
 - ▶ $B = [1 \quad 8 \quad 13 \quad 27]$
 - ▶ $\text{rank}(B:A) = [0 \quad 1 \quad 2 \quad 4]$
 - ▶ $\text{rank}(A:B) = [1 \quad 3 \quad 3 \quad 3 \quad 4 \quad 4 \quad 4]$
-
- ▶ Use binary search
 - ▶ Consider a multithreaded vs Hadoop implementation



Parallel Sort

MERGE SORT

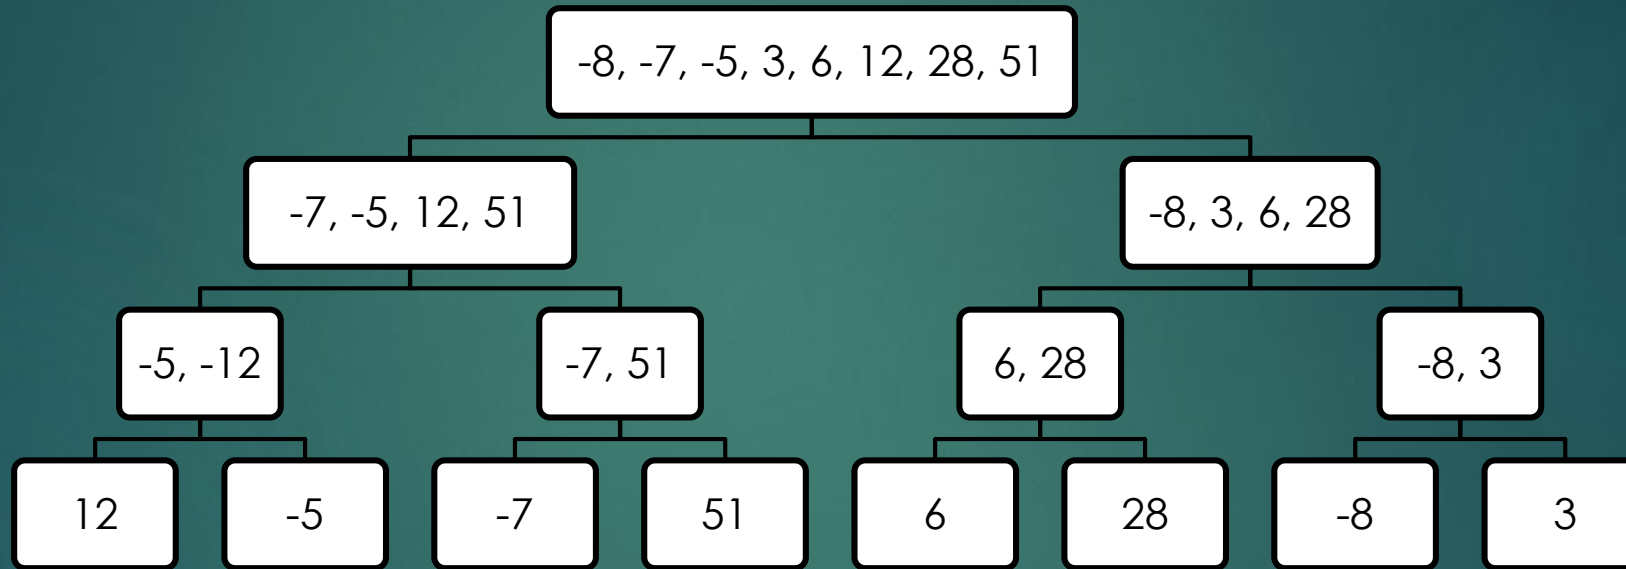


Divide & Conquer Merge Sort

- Divide X into X_1 and X_2
- Sort X_1 and X_2
- Merge X_1 and X_2
- Uses a Binary Tree
 - ▶ Bottom-up approach
 - ▶ Start with the leaves
 - ▶ Climb to the root
 - ▶ Merge the branches
- Requires parallel **Merge**



example



Input ↑




Merge sort

```
b = Merge_Sort(a,n)
    if n < 100
        return seqSort(a, n);
    b1 = Merge_Sort(a[0,...,n/2-1], n/2);
    b2 = Merge_Sort(a[n/2,...,n-1], n/2);
    return Merge (b1, b2);
```



Merge Sort - Complexity



Parallel Merge



Merging two lists of lengths n, m

- ▶ Problem description ($m \leq n$)
 - ▶ Given, $A = (a_1, a_2, \dots, a_n)$ and $B = (b_1, b_2, \dots, b_m)$
 - ▶ $a_i < a_{i+1} \quad \forall i$
 - ▶ $b_i < b_{i+1} \quad \forall i$
 - ▶ $A \cap B = \emptyset$
 - ▶ Build $C = (c_1, c_2, \dots, c_{n+m})$
 - ▶ $c_i \in A \cup B$
 - ▶ $c_i < c_{i+1} \quad \forall i$



Merging two sorted lists

- Best Sequential Time: $O(n)$
- Parallel Merge:
 - ▶ Tradeoffs between
 - ▶ Depth-Optimal
 - ▶ Work-Optimal




Merging using Ranking

- ▶ Assume elements in A and B are distinct
- ▶ Let C be the merged result. Given,
 - ▶ $x \in C$
 - ▶ $\text{rank}(x: C) = i$
 - ▶ $c_i = x$
- ▶ Property
$$\text{rank}(x: C) = \text{rank}(x: A) + \text{rank}(x: B)$$
- ▶ Solution to the merging problem,
 - ▶ Find $\text{rank}(A: B)$ and $\text{rank}(B: A)$
 - ▶ Parallel searches using $p = nm$, $D = O(1)$ but $W = O(n^2)$
 - ▶ Concurrent binary searches, $D = O(\log n)$ and $W = O(n \log n)$
 - ▶ Goal: Parallelize with optimal work

Recall that an algorithm is work optimal iff $W_p = W_{seq}$



Example



Work-optimal merge - Merge1

- ▶ $A = [a_1, \dots, a_n], \quad B = [b_1, \dots, b_m], \quad n \geq m$
- 1. Partition B into $\frac{m}{\log m}$ blocks
 - ▶ Size of each block $\rightarrow \log m$
- 2. parallel for $i = 1 : m/\log m$
 - ▶ $R_i \leftarrow \text{rank}(b_{i \log m} : A)$ using sequential binary search
- 3. Partition A accordingly
 - ▶ Block $A_i : (a_{R_{i-1}+1}, \dots, a_{R_i})$
- 4. Merge blocks A_i and B_i in $O(\log m)$ time using sequential merge
 - ▶ But if $|A_i| \gg |B_i| = \log m$, then recurse ... Merge1(B_i, A_i)

Work ?
Depth ?



Sequential Sorting

- ▶ What is the complexity ?

$O(?)$



Sequential Sorting

- ▶ Comparison based
 - ▶ $\mathcal{O}(n \log n)$
 - ▶ Can we sort faster than $\mathcal{O}(n \log n)$?
- ▶ Non-comparison based
 - ▶ $\mathcal{O}(n)$



Bucket sort

- ▶ Assume input is uniformly distributed over an interval $[a, b]$
- ▶ Divide interval into m equal sized intervals (**buckets**)
- ▶ Drop numbers into appropriate buckets
- ▶ Sort each bucket (say using quicksort)
 - ▶ $\mathcal{O}\left(n \log\left(\frac{n}{m}\right)\right)$
 - ▶ For $m = \mathcal{O}(n) \rightarrow \mathcal{O}(n)$ sorting
- ▶ Radix sort
- ▶ dense, uniform distribution




Parallel Quicksort




p_1




p_2



p_3




p_4

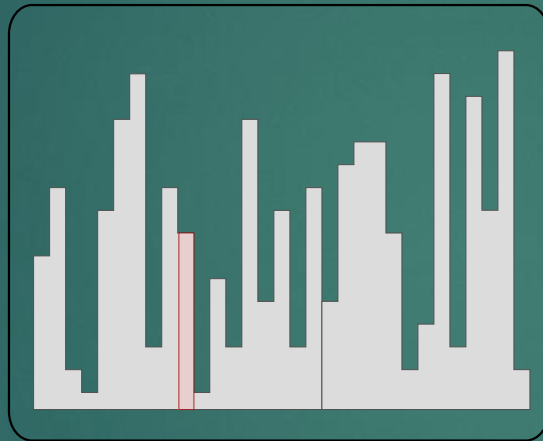


Parallel Quicksort


p_1




p_2



p_3



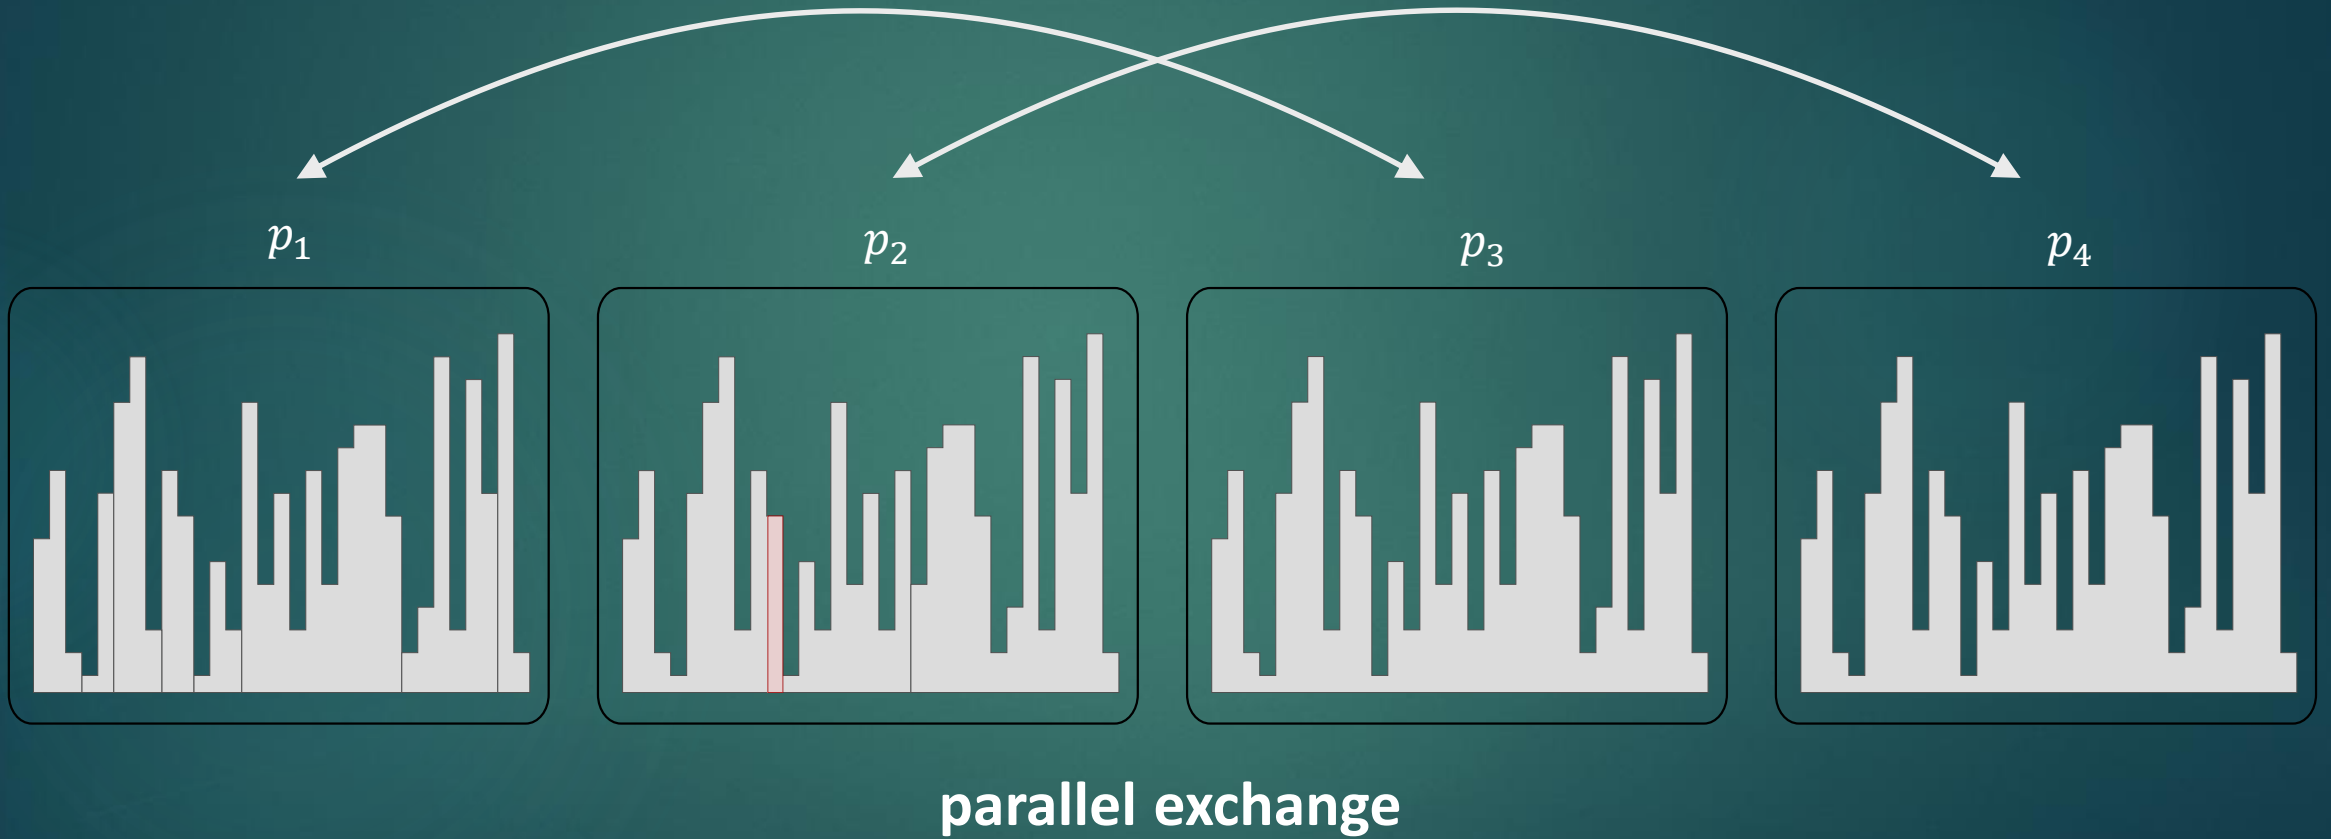
p_4



parallel median selection




Parallel Quicksort




Parallel Quicksort

p_1




p_2



p_3



p_4






Sample Sort


p_1



p_2



p_3




p_4




Sample Sort




p_1




p_2




p_3



p_4




Sample Sort

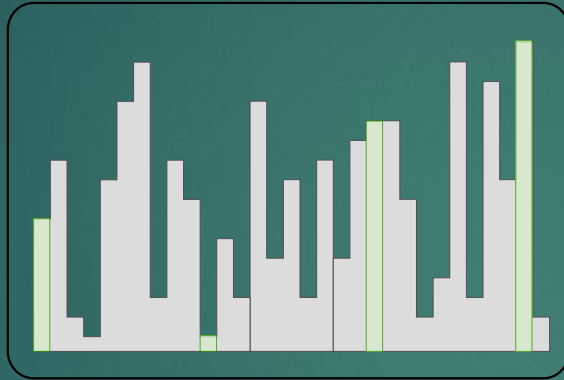


Sample Sort


p_1




p_2



p_3



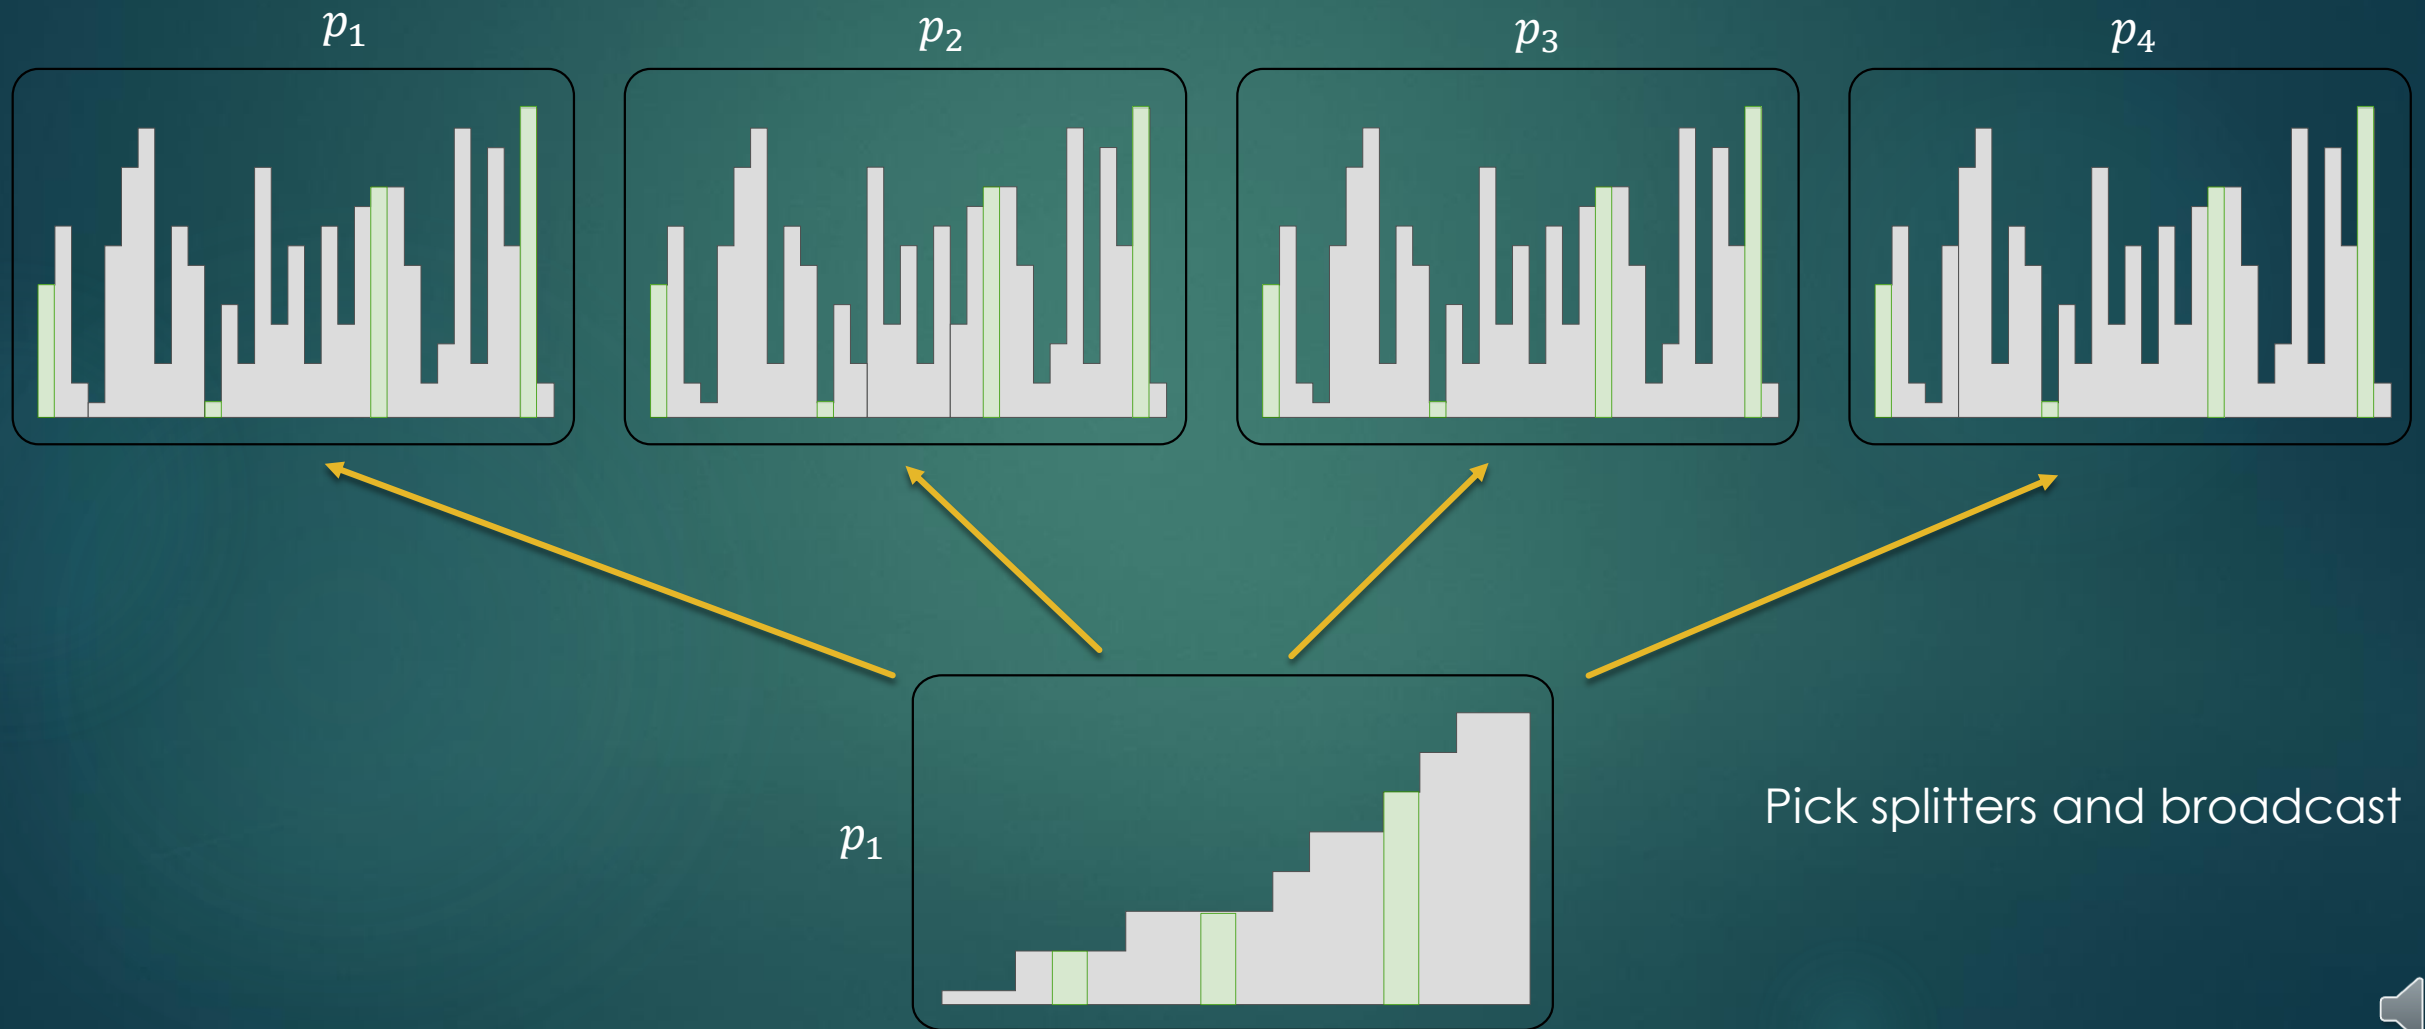
p_4



p_1

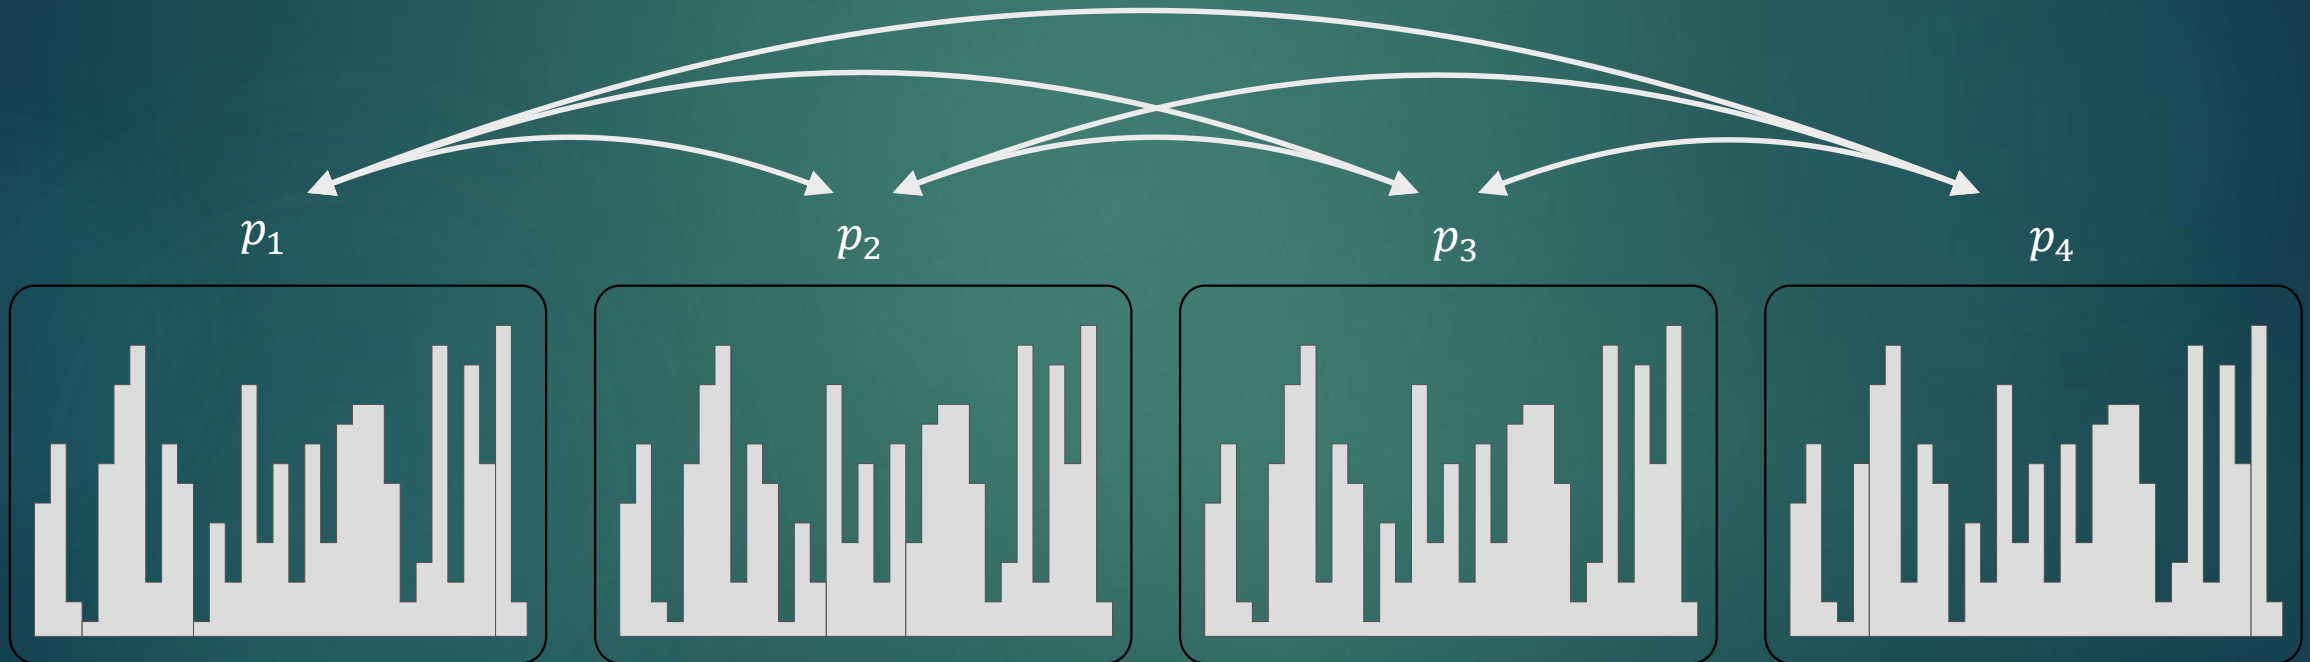


Sample Sort



Sample Sort

bucket data & all2all exchange



Sample sort

- ▶ randomly partition input in $\frac{n}{p}$ points
- ▶ sort locally
- ▶ select p splitters/processor (evenly)
 - ▶ guarantees no more $2 \cdot n/p$ elements / bucket (**proof?**)
- ▶ gather(splitters) in p_0
- ▶ sort splitters in p_0 and create buckets
 - ▶ block partition using p binary search on n/p sorted seq.
- ▶ Exchange data
- ▶ sort



Sample Sort

- ▶ Sort locally
- ▶ Select $p - 1$ splitters per process
- ▶ Gather splitters in p_0
- ▶ Sort splitters in p_0
- ▶ Broadcast splitters
- ▶ Sort again

$$\mathcal{O}\left(\frac{n}{p} \log \frac{n}{p}\right)$$

$$\mathcal{O}(p)$$

$$\mathcal{O}(p^2)$$

$$\mathcal{O}(p^2 \log p)$$

$$\mathcal{O}(p \log p)$$

$$\mathcal{O}\left(\frac{n}{p} \log \frac{n}{p}\right)$$



Sample Sort – load balance

- ▶ Guarantees no more $\frac{2n}{p}$ elements / bucket
- ▶ Proof:

All entries on p_i must be $> s_{i-1}$ and $\leq s_i$

$(i-2)p + \frac{p}{2}$ elements of the sample $\leq s_i$

$$\rightarrow \text{lower bound elements} = \frac{\left((i-2)p + \frac{p}{2}\right)n}{p^2}$$

$(p-i)p - \frac{p}{2}$ elements of the sample $> s_i$

$$\rightarrow \text{upper bound elements} = \frac{\left((p-i)p - \frac{p}{2}\right)n}{p^2} + \frac{n}{p^2} - 1$$

Maximum number of elements on processor i ,

$$n - ub - lb = \frac{2n}{p} - \frac{n}{p^2} + 1 \leq \frac{2n}{p}$$

■

