

Randomized Algorithms



Last time ...

- ▶ Network topologies
- ▶ Intro to MPI
- ▶ Matrix-matrix multiplication



Today ...

- ▶ MPI I/O
- ▶ Randomized Algorithms
 - ▶ Parallel k -Select
 - ▶ Graph coloring
- ▶ Assignment 2


Parallel I/O

- ▶ Goal of Parallel I/O is to use parallelism to increase bandwidth
- ▶ Parallel I/O can be hard to coordinate and optimize if working directly at the level of Lustre(DFS) API or POSIX I/O Interface
- ▶ Therefore, specialists implement a number of intermediate layers for coordination of data access and mapping from application layer to I/O layer
- ▶ Hence, application developers only have to deal with a high-level interface built on top of a software stack that in turn sits on top of the underlying hardware




Typical Pattern: Parallel programs doing sequential I/O

- ▶ All processes send data to master process, and then the process designated as master writes the collected data to the file
- ▶ This sequential nature of I/O can limit performance and scalability of many applications



Desired Pattern: Parallel Programs doing Parallel I/O

- ▶ Multiple processes participating in reading data from or writing data to a common file in parallel
- ▶ This strategy improves performance and provides a single file for storage and transfer purposes



Simple Solution?



MPI for Parallel I/O

- ▶ Reading and writing in parallel is like receiving and sending messages
- ▶ Hence, an MPI-like machinery is a good setting for Parallel I/O (think MPI communicators and MPI datatypes)
- ▶ MPI-I/O featured in MPI-2 which was released in 1997, and it interoperates with the file system to enhance I/O performance for distributed memory applications
- ▶ Given N number of processes, each process participates in reading or writing a portion of a common file



Using MPI I/O

- ▶ `MPI_File_open`, `MPI_File_close`
- ▶ There are three ways of positioning where the read or write takes place for each process:
 - ▶ Use individual file pointers (e.g., `MPI_File_seek/MPI_File_read`)
 - ▶ Calculate byte offsets (e.g., `MPI_File_read_at`)
 - ▶ Access a shared file pointer (e.g., `MPI_File_seek_shared`, `MPI_File_read_shared`)
- ▶ Collective I/O
 - ▶ forces all processes in the communicator to read/write data simultaneously and to wait for each other
 - ▶ `MPI_File_read_(at)_all`, `MPI_File_write_(at)_all`,



Randomized Algorithms



Recall ...

- ▶ Two main strategies for parallelization
 - ▶ Divide & Conquer
 - ▶ **Randomization**

Ensure that processors can make local decisions which, with high probability, add up to good global decisions

Sampling → quicksort, samplesort



Randomization

- ▶ Sampling
- ▶ Symmetry breaking
 - ▶ Independent sets → today
- ▶ Load balancing



Parallel k -Select

- ▶ Given an array A with n elements,
 - ▶ $\text{Select}(A, k)$ returns the k^{th} largest element in A
 - ▶ $\text{Select}(A, n/2), \text{Select}(A, \frac{n+1}{2}) \rightarrow$ median selection
 - ▶ $\text{Select}(A, 1) \rightarrow$ minimum
 - ▶ $\text{Select}(A, n) \rightarrow$ maximum
-
- ▶ Note that for sample sort, we can perform $(p - 1)$ k -selects simultaneously to obtain the $p - 1$ splitters

$$\text{Select}\left(A, \frac{n}{p}, \frac{2n}{p}, \dots, \frac{(p-1)n}{p}\right)$$



Parallel k -Select

- ▶ Select Splitter(s) randomly $\rightarrow (k_1, k_2)$
- ▶ Estimate ranks for the splitters (R_1, R_2) and partition A into three sets
 - ▶ Sequential ranking (count) on each process $\rightarrow \mathcal{O}\left(\frac{n}{p}\right)$
 - ▶ Reduction to obtain global ranking $\rightarrow \mathcal{O}(\log p)$
- ▶ Recurse on the appropriate set
 - ▶ depending on whether $k \in [0, R_1)$, $[R_1, R_2)$, or $[R_2, n)$
 - ▶ Recursion on a strictly smaller set \rightarrow will terminate
- ▶ Tradeoff between number of splitters at each iteration and number of iterations to converge
 - ▶ Choosing $p - 1$ splitters efficient for samplesort
 - ▶ What to do for quicksort?



Graph Algorithms

- ▶ Will be covered in detail later ...
- ▶ Graph $\mathcal{G} = (V, E)$, vertices and edges
- ▶ Matrices \rightarrow Graphs
 - ▶ Adjacency graph of a matrix A
 - ▶ Edge (i, j) exists iff $A_{ij} \neq 0$
 - ▶ Edge weight, W_{ij} , can be the A_{ij} value
- ▶ Graphs \rightarrow Matrices
 - ▶ Adjacency matrix of a weighted graph
 - ▶ Default weight 1, vertex value is in-degree
 - ▶ Symmetric \rightarrow undirected graphs
 - ▶ Unsymmetric \rightarrow directed graphs



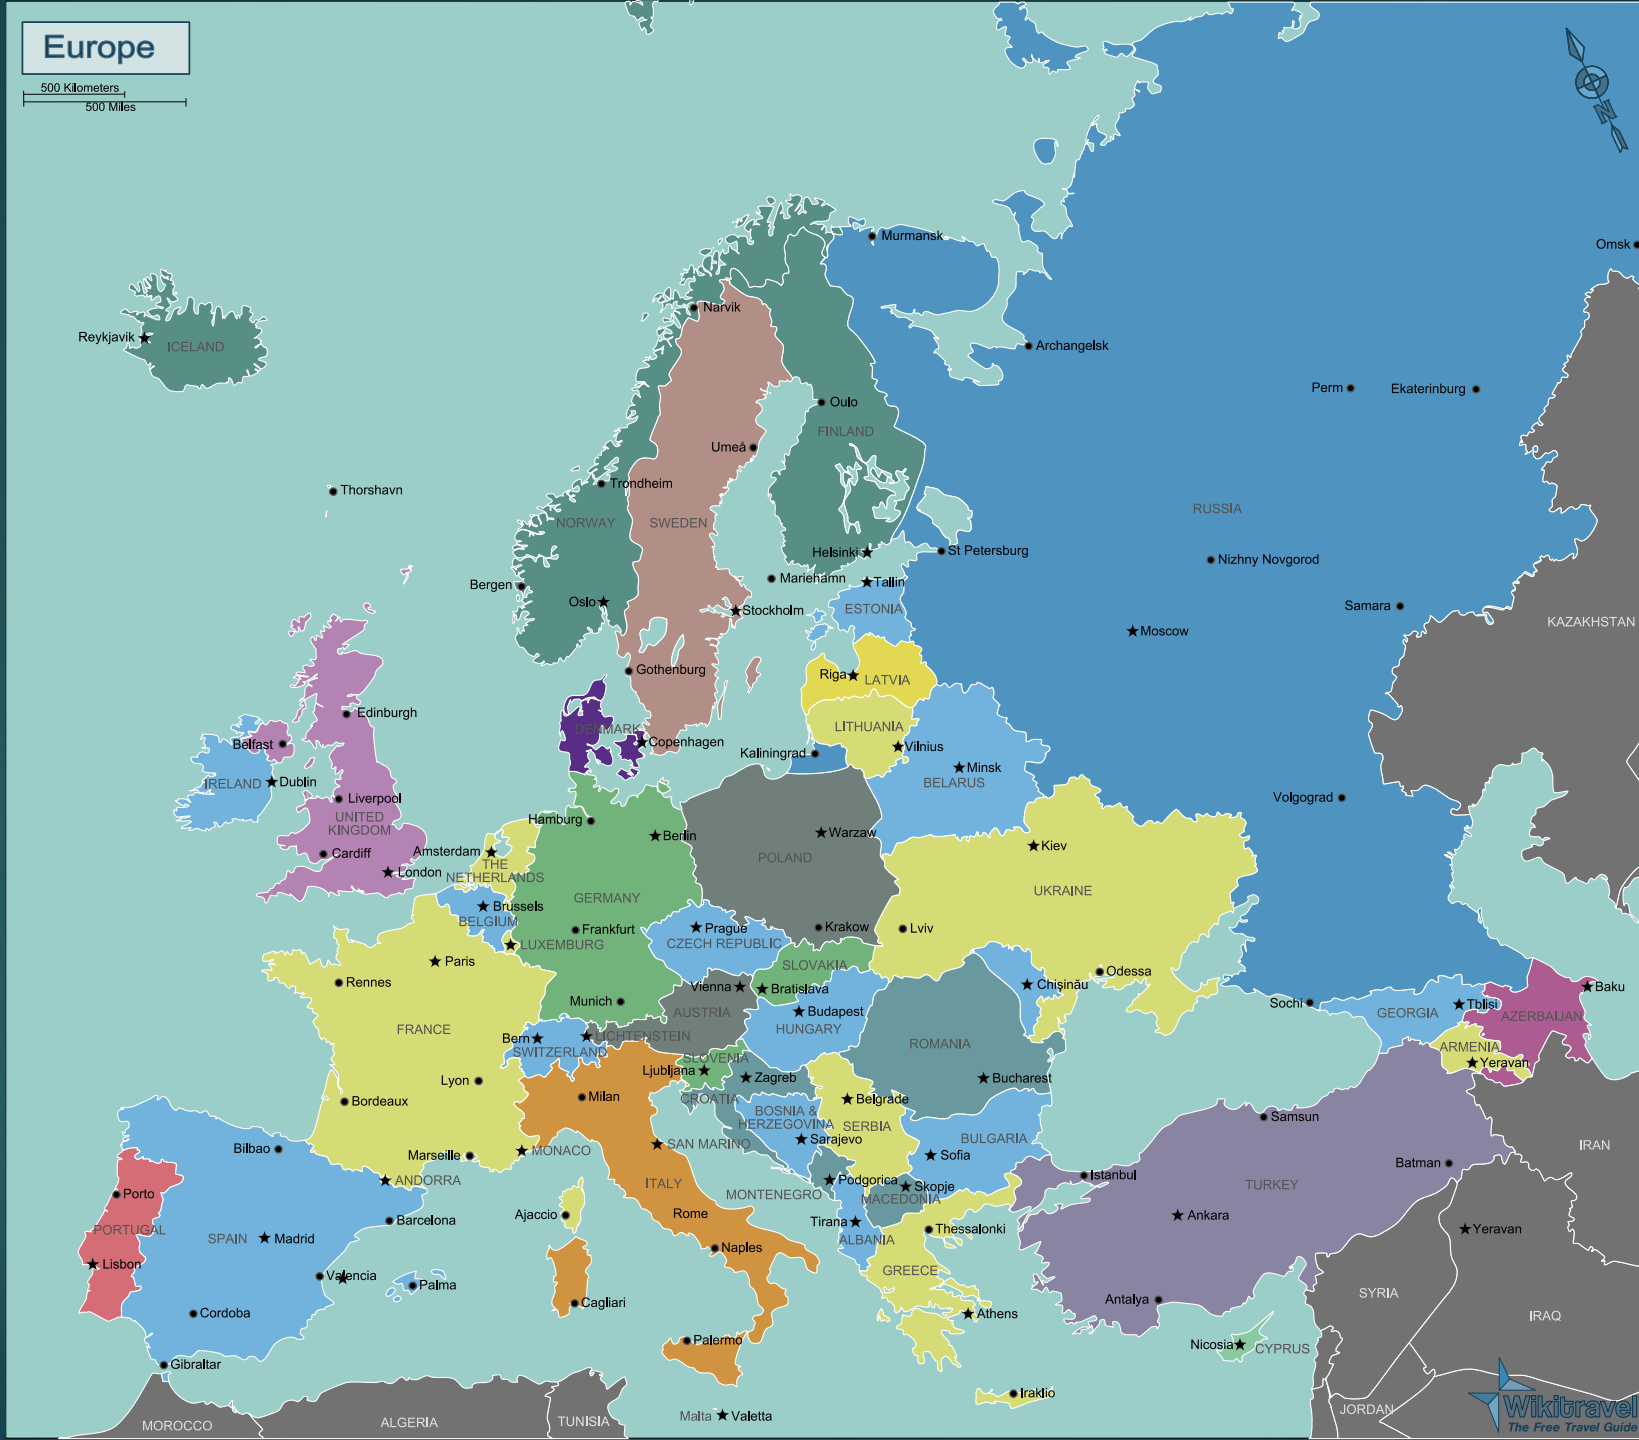
Graph Algorithms

- ▶ Graph partitioning
 - ▶ NP Hard problem
 - ▶ We will cover in detail
 - ▶ Coloring
- ▶ Graph Laplacian & Eigenproblem
- ▶ Breadth First Search (BFS)
- ▶ Depth First Search (DFS)
- ▶ Connected components
 - ▶ Spanning Trees



Europe

500 Kilometers
500 Miles



CS 5965/6965 - Big Data Systems - Fall 2014



Sequential Greedy Algorithm


- ▶ $n = |V|$
- ▶ Choose a random permutation $p(1), \dots, p(n)$ of numbers $1, \dots, n$
- ▶ $U \leftarrow V$
- ▶ for $i \leftarrow 1$ to n
 - ▶ $v \leftarrow p(i)$
 - ▶ $S \leftarrow \{\text{colors of all colored neighbors of } v\}$
 - ▶ $c(v) \leftarrow \text{smallest color } \notin S$
 - ▶ $U \leftarrow U \setminus \{v\}$

Bounds?



(Maximal) Independent Set

Independent Set: no two vertices share a common edge



Parallel Graph Coloring

- ▶ Any independent set can be colored in parallel
- ▶ $U \leftarrow V$
- ▶ while $|U| > 0$ do in parallel
 - ▶ Choose an independent set I from U
 - ▶ Color all vertices in I
 - ▶ $U \leftarrow U \setminus I$
- ▶ Optimal Coloring \rightarrow color using smallest color
- ▶ Balanced Coloring \rightarrow use all colors equally



Maximal Independent Set (Luby)

- ▶ find largest MIS from graph
- ▶ Color all with the same color and remove from graph
- ▶ Recurse

$I \leftarrow \emptyset$

$V' \leftarrow V$

while $|V'| > 0$ **do**

 choose and independent set I' from V'

$I \leftarrow I + I'$

$X \leftarrow I' + N(I')$

$V' \leftarrow V' - X$




How to choose independent sets in parallel?

- ▶ Assign a random weight to each vertex
- ▶ Choose vertices that are a local maxima
- ▶ $\mathcal{O}((c + 1) \log|V|)$ algorithm
 - ▶ for sparse graphs




Luby's Algorithm



initially assigned random numbers






CS 5965/6965 - Big Data Systems - Fall 2014




Luby's Algorithm



Find MIS




Luby's Algorithm



Find MIS, color all the same color




Luby's Algorithm



repeat




Luby's Algorithm



repeat




Luby's Algorithm



repeat



Luby's Algorithm



repeat



Jones-Plassmann Coloring

- ▶ Not necessary to create a new random permutation of vertices every time
- ▶ Use vertex number to resolve conflicts
- ▶ Does not find a MIS at each step
- ▶ Instead,
 - ▶ Find independent set
 - ▶ Not assigned the same color
 - ▶ Color individually using smallest available color




Jones-Plassmann Coloring

- ▶ $U \leftarrow V$
- ▶ **while** $|U| > 0$ **do**
 - ▶ **for all** vertices $v \in U$ **do in parallel**
 - ▶ $I \leftarrow \{v \mid w(v) > w(u) \forall u \in N(v)\}$
 - ▶ **for all** vertices $v' \in I$ **do in parallel**
 - ▶ $S \leftarrow \{\text{colors of } N(v')\}$
 - ▶ $c(v') \leftarrow \text{minimum color } \notin S$
 - ▶ $U \leftarrow U \setminus I$




Jones-Plassmann Coloring



initially assigned random numbers




Jones-Plassmann Coloring



If local maxima, assign lowest color




Jones-Plassmann Coloring



repeat, considering only uncolored vertices




Jones-Plassmann Coloring



repeat, considering only uncolored vertices




Jones-Plassmann Coloring



repeat, considering only uncolored vertices



Jones-Plassmann Coloring



repeat, considering only uncolored vertices

