

Clustering



Last time ...

- ▶ PageRank
 - ▶ Topic-sensitive
 - ▶ Link Spam
- ▶ Today
 - ▶ Clustering



Points, Spaces & Distances

- ▶ Points
- ▶ Spaces (normed vector space)
 - ▶ e.g. Euclidean
 - ▶ $\|x\| > 0$, if $x \neq 0$
 - ▶ $\|\alpha x\| = |\alpha| \|x\|$ for any scalar α
 - ▶ Triangle inequality, $\|x + y\| \leq \|x\| + \|y\|$
- ▶ Distance measure
 - ▶ non-negative
 - ▶ Symmetric
 - ▶ Triangle inequality



Clustering

- ▶ Hierarchical or agglomerative algorithms
- ▶ Point assignment
- ▶ Euclidean or arbitrary distance measure (e.g. Riemann, Hyperbolic)
 - ▶ Centroid of cluster
- ▶ Will the data fit in main memory
- ▶ Dimensionality



Non-Euclidian Spaces

- ▶ Need to pick appropriate distance measure
- ▶ Centroid not appropriate for clustering
 - ▶ Choose a sample in place of centroid
- ▶ Hyperplane partitions



K-means clustering

- ▶ Point assignment algorithm
 - ▶ Assume Euclidean space
 - ▶ Assume number of clusters, k , is known in advance
-
1. Choose k points that are likely to be in different clusters
 2. Foreach remaining point p do
 1. Find the cluster (centroid) closest to p
 2. Add p to this cluster
 3. Update the centroid for the cluster



Initializing the clusters

- ▶ Pick points that are as far away from one another as possible
 - ▶ Pick a random first point
 - ▶ Add additional points that maximize the minimal distance to already selected points
- ▶ Cluster a sample of the data and choose clusters
 - ▶ Random samples
 - ▶ Hierarchical clustering



Bradley, Fayyad & Reina (BFR)

- ▶ High-dimensional Euclidean space
- ▶ Clusters are normally distributed about the centroid (assumption)
- ▶ Start by choosing k centroids
- ▶ Read data in chunks
- ▶ Keep 3 types of information in memory
 - ▶ Discard Set $\rightarrow 2d + 1$ values (number of points, sum, sum of squares)
 - ▶ Compressed Set $\rightarrow 2d + 1$ values
 - ▶ Retained Set




Bradley, Fayyad & Reina (BFR)

- ▶ First, all points that are sufficiently close to the centroid of a cluster are added to that cluster (Mahalanobis distance)
- ▶ For the points that are not sufficiently close to any centroid, we cluster them, along with the points in the retained set
- ▶ Merge miniclusters (new clusters + existing compressed set)
- ▶ Final processing of points in the retained set and the miniclusters in the compressed set
 - ▶ Discard → outliers
 - ▶ merge



Clustering Using REpresentatives

- ▶ CURE
 - ▶ Assumes Euclidean space
 - ▶ No assumptions on the shape of clusters
 - ▶ Uses a collection of representative points
1. Sample data and cluster hierarchically
 2. Select representative points for each cluster
 3. Move each of the representative points a fixed fraction of the distance between its location and the centroid of its cluster
 4. Merge clusters if close



Clustering in non-Euclidean spaces



Clustering in non-Euclidean spaces

- ▶ Use a combination of hierarchical and point-assignment
- ▶ Represent clusters by sample points in memory
- ▶ Organize clusters hierarchically in a tree
- ▶ Insert new points by traversing the tree



Representing Clusters (Ganti et al.)

- ▶ Clusters represented by a set of features
 - ▶ N , total number of points in the cluster
 - ▶ Clusteroid → point in the cluster that has minimum cumulative squared distance to all other points
 - ▶ The cumulative squared distance (CSD) of the clusteroid
 - ▶ k closest points to the clusteroid and their CSD
 - ▶ New clusteroid will be from one of these
 - ▶ k furthest points to the clusteroid and their CSD
 - ▶ Needed to decide whether to merge two clusters



Initializing the tree

- ▶ Tree structure is similar to a B-tree
- ▶ Each leaf node holds as many clusters as possible
- ▶ Interior nodes holds a sample of the clusteroids on its subtrees
- ▶ Initialize by hierarchically clustering a sample of the data
 - ▶ Choose only clusters that are of a specific size
 - ▶ These are the leaf nodes
 - ▶ Group clusters based on common ancestors
- ▶ Balance tree if necessary



Adding points

- ▶ Read points from disk and traverse down tree, using distance to clusteroid (samples)
- ▶ On reaching the leaf, pick and update the cluster representation
 - ▶ Increment N
 - ▶ Update the CSD of all $2k + 1$ points
 - ▶ Squared distance to the new point
 - ▶ Estimate the CSD of the new point, p
 - ▶ $\text{CSD}(p) = \text{CSD}(c) + Nd^2(p, c)$
 - ▶ Check if p is either the k nearest or farthest point
 - ▶ Check if one of the k nearest points is the new clusteroid



Clustering in a Streaming model

- ▶ Assume a sliding window of N points
- ▶ We are interested in the clusteroids of the best clusters formed from the most recent m points, for any $m \leq N$
- ▶ Pre-cluster the points so that queries can be answered
- ▶



Stream-Clustering (Babcock et al.)

- ▶ Partition and summarize data into buckets
 - ▶ Bucket sizes are powers of two $\rightarrow c \cdot 2^k$
 - ▶ Only one or two of any size
 - ▶ Bucket sizes are restrained to be non-decreasing as we go back in time
 - ▶ $\log N$ buckets
- ▶ Contents of a bucket
 - ▶ Size of the bucket
 - ▶ Timestamp of the bucket \rightarrow most recent point
 - ▶ Cluster representations (multiple clusters per bucket)
 - ▶ Number of points in the cluster
 - ▶ Clusteroid
 - ▶ Other info needed to merge and maintain the clusters



Initializing buckets

- ▶ p - smallest bucket size
- ▶ Every p points, create a new bucket. Timestamp the bucket
 - ▶ Cluster points
- ▶ Drop any bucket older than N
- ▶ If we have 3 buckets of size $p \rightarrow$ merge the two oldest
- ▶ Might have to propagate merges ($2p, 4p \dots$)



Merging buckets

- ▶ Size of new bucket is twice as large
- ▶ Timestamp of the new bucket is the newer of the two being merged
- ▶ Decide on whether to merge clusters
- ▶ Consider: k -means, Euclidean
 - ▶ Represent clusters using number of points (n) and centroid (c)
 - ▶ Choose $p = k$, or larger $\rightarrow k$ -means clustering while creating bucket
 - ▶ To merge, $n = n_1 + n_2$, $c = \frac{n_1 c_1 + n_2 c_2}{n_1 + n_2}$



Merging buckets

- ▶ Size of new bucket is twice as large
- ▶ Timestamp of the new bucket is the newer of the two being merged
- ▶ Decide on whether to merge clusters
- ▶ Consider: non-Euclidean (Ganti et al.)
 - ▶ Represent clusters using clusteroid and CSD
 - ▶ Need to choose new clusteroid while merging
 - ▶ Will be one of the k -points furthest from the clusteroids
 - ▶ $CSD_m(p) = CSD_1(p) + N_2(d^2(p, c_1) + d^2(c_1, c_2)) + CSD_2(c_2)$



Answering queries

- ▶ Given m , choose the smallest set of buckets that covers the most recent m points. At most $2m$ points
- ▶ Merge clusters

